

# Table of Contents

## About this course (10min)

- What is scientific computing?
- Textbook

## Lecture 1: Understanding our computing devices

- Get hardware information (Linux)

## Number system (20min)

- Integers
- Floating point numbers
- The distribution of floating point numbers

## Estimating the computing power of your devices (20min)

- Example 1: Matrix multiplication
- Example 2: axpy
- Example 3: modified axpy

## Programming on a device (30min)

- Your programs are compiled to binary
- Measuring the performance

## Summarize

## Next lecture

- Pre-reading

# About this course (10min)

---

## What is scientific computing?

---




Scientific computing is the collection of tools, techniques and theories required to solve on a computer the mathematical models of problems in science and engineering.

– Gene H. Golub and James M. Ortega

## Textbook

---



## Chapters

1. Scientific Computing
2. Systems of linear equations
3. Linear least squares
4. Eigenvalue problems
5. Nonlinear equations
6. Optimization
7. Interpolation
8. ~~Numerical integration and differentiation~~
9. ~~Initial value problems for ordinary differential equations~~
10. ~~Boundary value problems for ordinary differential equations~~
11. ~~Partial differential equations~~
12. Fast fourier transform
13. Random numbers and stochastic simulation

## Lecture 1: Understanding our computing devices

---

# What is inside a computer? (40min)

## Get hardware information (Linux)

---

</> Get hardware information

\$ lscpu

\$ lsmem

\$ top

show\_cpuinfo = ☐

```
• if show_cpuinfo run(`lscpu`) end
```

show\_meminfo = ☐

```
• if show_meminfo run(`lsmem`) end
```

show\_processinfo = ☐

```
• if show_processinfo run(`top -n 1 -b`) end
```

# Integers

- `bitstring(typemax(Int64))`

# Floating point numbers



- `exponent(0.15625f0)`

- `significand(0.15625f0)` *# the fraction*

- `typemax(Float64)`

- `bitstring(Inf)`

- prevfloat(Inf)

- `typemin(Float64)`

- `bitstring(-Inf)`

- `prevfloat(0.0)`

- `nextfloat(0.0)`

- Inf-Inf


- 0 \* NaN

- `bitstring(NaN)`

```
• xs = filter(!isnan, reinterpret(Float64, rand{Int64, npoints}));
```

From the linear scale plot, you will see data concentrated around 0 (each vertical bar is a sample)


```
• using Plots
```



```
• scatter(xs, zeros(length(xs)), xlim=(-1e300, 1e300), label="", size=(600, 100),  
  yaxis=:off, markersize=30, markershape=:vline)
```

If we use logarithmic x-axis

smearing\_factor =



```
• let  
•   logxs = sign.(xs) .* log10.(abs.(xs))  
•   ax = scatter(logxs, zeros(length(xs)), xlim=(-300, 300), label="", size=(600,  
100), yaxis=:off, markersize=30, markershape=:vline)  
•   a = -300:300  
•   # smoothen the distribution with Lorentz function  
•   m = 1/π/smearing_factor ./ (((a' .- logxs) ./ smearing_factor) .^ 2 .+ 1)  
•   plot!(ax, a, dropdims(sum(m, dims=1), dims=1); label="probability")  
• end
```

# Estimating the computing power of your devices (20min)

## Example 1: Matrix multiplication

$$C_{ik} = \sum_j A_{ij} \times B_{jk}$$

Let the matrix size be  $n \times n$ , the pseudocode for general matrix multiply (GEMM) is

```
for i=1:n
  for j=1:n
    for k=1:n
      C[i, k] = A[i, j] * B[j, k]
    end
  end
end
```

GEMM is CPU bottlenecked

`matrix_size =`

`benchmark_example1 =` ☐

Loading the package for benchmarking

```
• using BenchmarkTools
```

Loading the matrix multiplication function

```
• using LinearAlgebra: mul!
```

```
• if benchmark_example1
•   let
•     # creating random vectors with normal distribution/zero elements
•     A = randn(Float64, matrix_size, matrix_size)
•     B = randn(Float64, matrix_size, matrix_size)
•     C = zeros(Float64, matrix_size, matrix_size)
•     @benchmark mul!($C, $A, $B)
•   end
• end
```

Calculating the **floating point operations per second**

# FLOPS for computing GEMM

the number of floating point operations / the number of seconds

## Example 2: axpy

---

axpy! is memory I/O bottlenecked

axpy! (generic function with 1 method)

```
• function axpy!(a::Real, x::AbstractVector, y::AbstractVector)
•   @assert length(x) == length(y) "the input size of x and y mismatch, got
      $(length(x)) and $(length(y))"
•   @inbounds for i=1:length(x)
•       y[i] += a * x[i]
•   end
•   return y
• end
```

axpy\_vector\_size =

benchmark\_axpy = ☐

```
• if benchmark_axpy
•   let
•       x = randn(Float64, axpy_vector_size)
•       y = randn(Float64, axpy_vector_size)
•       @benchmark axpy!(2.0, $x, $y)
•   end
• end
```





# FLOPS for computing axpy

the number of floating point operations / the number of seconds

## Example 3: modified axpy

bad\_axpy! (generic function with 1 method)

- `function bad_axpy!(a::Real, x::AbstractVector, y::AbstractVector, indices::AbstractVector{Int})`
- `@assert length(x) == length(y) == length(indices) "the input size of x and y mismatch, got $(length(x)), $(length(y)) and $(length(indices))"`
- `@inbounds for i in indices`
- `y[i] += a * x[i]`
- `end`
- `return y`
- `end`

I will show this function is latency bottlenecked

bad\_axpy\_vector\_size =

benchmark\_bad\_axpy = ☐

- `@xbind benchmark_bad_axpy CheckBox()`

- `using Random`

- `if benchmark_bad_axpy`
- `let`
- `x = randn(Float64, bad_axpy_vector_size)`
- `y = randn(Float64, bad_axpy_vector_size)`
- `indices = randperm(bad_axpy_vector_size)`
- `@benchmark bad_axpy!(2.0, $x, $y, $indices)`
- `end`
- `end`

# FLOPS for computing bad axpy

the number of floating point operations / the number of seconds

## Programming on a device (30min)

---

### You program are compiled to binary

---

```
.text
.file "axpy!"
.globl "julia_axpy!_5041"                # -- Begin function julia_axpy!_5041
.p2align 4, 0x90
.type "julia_axpy!_5041",@function
"julia_axpy!_5041":                     # @"julia_axpy!_5041"
; | @ /home/leo/jcode/ModernScientificComputing/notebooks/1.understanding-our-computi
.cfi_startproc
# %bb.0:                                # %top
    pushq    %rbp
    .cfi_def_cfa_offset 16
    .cfi_offset %rbp, -16
    movq     %rsp, %rbp
    .cfi_def_cfa_register %rbp
    pushq    %r15
    pushq    %r14
    pushq    %rbx
    andq     $-32, %rsp
```

```
• with_terminal() do
•     x, y = randn(10), randn(10)
•     @code_native axpy!(2.0, x, y)
• end
```

Let us check an easier one

oneton (generic function with 1 method)

```
• function oneton(n::Int)
•     res = zero(n)
•     for i = 1:n
•         res+=i
•     end
•     return res
• end
```

```
.text
.file "oneton"
.globl julia_oneton_5153                # -- Begin function julia_oneton_5153
.p2align    4, 0x90
.type    julia_oneton_5153,@function
julia_oneton_5153:                    # @julia_oneton_5153
; | @ /home/leo/jcode/ModernScientificComputing/notebooks/1.understanding-our-computi
.cfi_startproc
# %bb.0:                                # %top
; | @ /home/leo/jcode/ModernScientificComputing/notebooks/1.understanding-our-computi
; | @ range.jl:5 within `Colon`
; | | @ range.jl:397 within `UnitRange`
; | | | @ range.jl:404 within `unitrange_last`
; | | testq    %rdi, %rdi
; | | LLL
; | | jle     .LBB0_1
# %bb.2:                                # %L17.preheader
; | @ /home/leo/jcode/ModernScientificComputing/notebooks/1.understanding-our-computi
```

```
• with_terminal() do
•     @code_native oneton(10)
• end
```

An instruct has a binary correspondence: [check the online decoder](#)

## Measuring the performance

```
• using Profile
```

```
profile_axpy = ☐
```

```
• if profile_axpy
•     with_terminal() do
•         # clear previous profiling data
•         Profile.init(; n=10^6, delay=0.001)
•         x, y = randn(100000000), randn(100000000)
•         @profile axpy!(2.0, x, y)
•         Profile.print()
•     end
• end
```

# How does profiling work?

\* function call stack

\* two approaches: instrumentation and sampling

## Summarize

---

1. understanding the components of our computing devices
2. the bottlenecks of our computing devices
3. how to get our program compiled and executed
4. how to measure the performance of a program with profiling

## Next lecture

---

We have have a coding seminar. I will show you some cheatsheets about

- Linux operation system
- Vim
- Git
- SSH
- Julia installation Guide

Please bring your laptops and get your hands dirty! If you are already an expert, please let me know, I need some help in preparing the cheatsheets.

## Pre-reading

---

Strong recommended course: [missing-semester](#)

- 1/13: Course overview + the shell (1 - expert)
- 1/14: Shell Tools and Scripting (4 - basic)
- 1/15: Editors (Vim) (2 - basic)
- 1/16: Data Wrangling
- 1/21: Command-line Environment
- 1/22: Version Control (Git) (3 - expert)
- 1/23: Debugging and Profiling
- 1/27: Metaprogramming (build systems, dependency management, testing, CI)
- 1/28: Security and Cryptography
- 1/29: Potpourri
- 1/30: Q&A

#### Note

- Yellow backgrounded lectures are required by AMAT5315
- (n - basic) is the reading order and the level of familiarity