

```
1 using PlutoUI, Test
```

Table of Contents

Simulting lattice gas cellular automata

Cellular automata

CUDA programming with Julia

Goal

Run a CUDA program

CUDA libraries and Kernel Programming

Appendix: The Navier-Stokes equation

```
1 TableOfContents()
```


Simulting lattice gas cellular automata

Cellular automata

- A descretized space and time,
- A state defined on the space,
- A simple set of rules (local & finite) to describe the evolution of the state.

Reference:

- Hardy J, Pomeau Y, De Pazzis O. Time evolution of a two-dimensional model system. I. Invariant states and time correlation functions[]. Journal of Mathematical Physics, 1973, 14(12): 1746-1759.



- Particles exist only on the grid points, never on the edges or surface of the lattice.
- Each particle has an associated direction (from one grid point to another immediately adjacent grid point).
- Each lattice grid cell can only contain a maximum of one particle for each direction, i.e., contain a total of between zero and four particles.

The following rules also govern the model:

- A single particle moves in a fixed direction until it experiences a collision.
- Two particles experiencing a head-on collision are deflected perpendicularly.
- Two particles experience a collision which isn't head-on simply pass through each other and continue in the same direction.
- Optionally, when a particles collides with the edges of a lattice it can rebound.

CUDA programming with Julia

CUDA programming is a parallel computing platform and programming model developed by NVIDIA for performing general-purpose computations on its GPUs (Graphics Processing Units). CUDA stands for Compute Unified Device Architecture.

References:

1. [JuliaComputing/Training](#)
2. [arXiv: 1712.03112](#)

Goal

1. Run a CUDA program
2. Write your own CUDA kernel
3. Create a CUDA project

Run a CUDA program

1. Make sure you have a NVIDIA GPU device and its driver is properly installed.

```
Process('nvidia-smi', ProcessExited(0))
```

```
1 run('nvidia-smi')
```

```

Wed May 10 15:34:25 2023
+-----+
+
| NVIDIA-SMI 525.105.17    Driver Version: 525.105.17    CUDA Version: 12.0
|
+-----+-----+
+
| GPU   Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC
| Fan   Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M.
|                                           | MIG M.
+-----+-----+
+
| 0   NVIDIA RTX A4500      Off          | 00000000:03:00.0 Off  |
| 33%   55C    P8      26W / 200W | 4947MiB / 20470MiB |    0%      Default
|                                           |
+-----+-----+
+
+
+-----+
+
+
Processes:
+-----+
| GPU   GI    CI          PID    Type    Process name                      GPU Memory
|   ID   ID                                 Name                               Usage
+-----+

```

2. Install the [CUDA.jl](#) package, and disable scalar indexing of CUDA arrays.

CUDA.jl provides wrappers for several CUDA libraries that are part of the CUDA toolkit:

- Driver library: manage the device, launch kernels, etc.
- CUBLAS: linear algebra
- CURAND: random number generation
- CUFFT: fast fourier transform
- CUSPARSE: sparse arrays
- CUSOLVER: decompositions & linear systems

There's also support for a couple of libraries that aren't part of the CUDA toolkit, but are commonly used:

- CUDNN: deep neural networks
- CUTENSOR: linear algebra with tensors

```
1 using CUDA; CUDA.allowscalar(false)
```

```
1 CUDA.versioninfo()
```

```
CUDA runtime 12.1, artifact installation
CUDA driver 12.1
NVIDIA driver 525.105.17, originally for CUDA 12.0

Libraries:
- CUBLAS: 12.1.0
- CURAND: 10.3.2
- CUFFT: 11.0.2
- CUSOLVER: 11.4.4
- CUSPARSE: 12.0.2
- CUPTI: 18.0.0
- NVML: 12.0.0+525.105.17

Toolchain:
- Julia: 1.9.0-rc3
- LLVM: 14.0.6
- PTX ISA support: 3.2, 4.0, 4.1, 4.2, 4.3, 5.0, 6.0, 6.1, 6.3, 6.4, 6.5, 7.0,
7.1, 7.2, 7.3, 7.4, 7.5
- Device capability support: sm_37, sm_50, sm_52, sm_53, sm_60, sm_61, sm_62, s
m_70, sm_72, sm_75, sm_80, sm_86

1 device:
0: NVIDIA RTX A4500 (sm_86, 14.871 GiB / 19.990 GiB available)
```

3. Choose a device (if multiple devices are available).

```
CUDA.DeviceIterator() for 1 devices:  
0. NVIDIA RTX A4500
```

```
1 devices()
```

```
dev = CuDevice(0): NVIDIA RTX A4500
```

```
1 dev = CuDevice(0)
```

grid > block > thread

1024

```
1 attribute(dev, CUDA.DEVICE_ATTRIBUTE_MAX_THREADS_PER_BLOCK)
```

1024

```
1 attribute(dev, CUDA.CU_DEVICE_ATTRIBUTE_MAX_BLOCK_DIM_X)
```

2147483647

```
1 attribute(dev, CUDA.CU_DEVICE_ATTRIBUTE_MAX_GRID_DIM_X)
```

4. Create a CUDA Array

```
10-element CuArray{Float32, 1, CUDA.Mem.DeviceBuffer}:
```

```
0.0  
0.0  
0.0  
0.0  
0.0  
0.0  
0.0  
0.0  
0.0  
0.0  
0.0
```

```
1 CUDA.zeros(10)
```

```
cuarray1 = 10-element CuArray{Float32, 1, CUDA.Mem.DeviceBuffer}:
```

```
1.0669513  
-1.0011338  
-0.45724186  
0.6212741  
-0.90845925  
0.36932236  
-0.33143434  
2.1754394  
0.45788017  
0.97185755
```

```
1 cuarray1 = CUDA.randn(10)
```

Test Passed

Thrown: `ErrorException`

```
1 @test_throws ErrorException cuarray1[3]
```

9.542758f0

```
1 CUDA.@allowscalar cuarray1[3] += 10
```

Upload a CPU Array to GPU

```
10-element CuArray{Float64, 1, CUDA.Mem.DeviceBuffer}:  
-0.5106224197359117  
 0.08435781951985412  
 1.3145343984144213  
 0.7198074551780037  
 1.9659466115028255  
 0.31945629591804797  
-0.5025034746872823  
-1.7102368720734844  
 0.3221473609865356  
-1.2079721286819451
```

```
1 CuArray(randn(10))
```

5. Compute

Computing a function on GPU Arrays

1. Launch a CUDA job - a few micro seconds
2. Launch more CUDA jobs...
3. Synchronize threads - a few micro seconds

Computing matrix multiplication.

0.202039044

```
1 @elapsed rand(2000,2000) * rand(2000,2000)
```

InterruptException:

```
1 @elapsed CUDA.@sync CUDA.rand(2000,2000) * CUDA.rand(2000,2000)
```

WARNING: Force throwing a SIGINT



Broadcasting a native Julia function Julia -> LLVM (optimized for CUDA) -> CUDA

poor_besselj (generic function with 1 method)

```
1 # this function is copied from lecture 9
2 function poor_besselj(v::Int, z::T; atol=eps(T)) where T
3     k = 0
4     s = (z/2)^v / factorial(v)
5     out = s::T
6     while abs(s) > atol
7         k += 1
8         s *= -(k+v) * (z/2)^2 / k
9         out += s
10    end
11    out
12 end
```

factorial (generic function with 1 method)

```
1 factorial(n) = n == 1 ? 1 : factorial(n-1)*n
```

x = 1001-element CuArray{Float64, 1, CUDA.Mem.DeviceBuffer}:

```
0.0
0.01
0.02
0.03
0.04
0.05
0.06
⋮
9.95
9.96
9.97
9.98
9.99
10.0
```

```
1 x = CUDA.CuArray(0.0:0.01:10)
```

1001-element CuArray{Float64, 1, CUDA.Mem.DeviceBuffer}:

```
0.0
0.0049997500093746875
0.009998000299960006
0.014993252277441754
0.01998400959488256
0.02496877927248
0.029946072812618307
⋮
1.6337875313577208e-5
0.005418409574627547
0.00675994621802067
4.02498425561632e-6
2.043448513104033e-9
9.927485697704232e-6
```

```
1 poor_besselj.(1, x)
```

```
1 using BenchmarkTools
```

6. manage your GPU devices

```
nvml_dev = NVML.Device(0): NVIDIA RTX A4500
```

```
1 nvml_dev = NVML.Device(parent_uuid(device()))
```

```
62.568
```

```
1 NVML.power_usage(nvml_dev)
```

```
(compute = 0.99, memory = 0.0)
```

```
1 NVML.utilization_rates(nvml_dev)
```

```
Dict(372567 => (used_gpu_memory = 673185792), 312821 => (used_gpu_memory = 811597824))
```

```
1 NVML.compute_processes(nvml_dev)
```

CUDA libraries and Kernel Programming

Please check [lib/CUDATutorial](#)

Appendix: The Navier-Stokes equation

Reference: <https://youtu.be/Ra7aQlenTb8>

The million dollar equation (Navier-Stokes equations)



The navier stokes equation describes the fluid dynamics, which contains the following two parts.

The first one describes the conservation of volume

$$\nabla \cdot \underbrace{u}_{\text{velocity } u \in \mathbb{R}^d} = 0$$

The second one describes the dynamics

$$\underbrace{\rho}_{\text{density}} \frac{du}{dt} = \underbrace{-\nabla p}_{\text{pressure}} + \underbrace{\mu \nabla^2 u}_{\text{viscosity (or friction)}} + \underbrace{f}_{\text{external force}} .$$